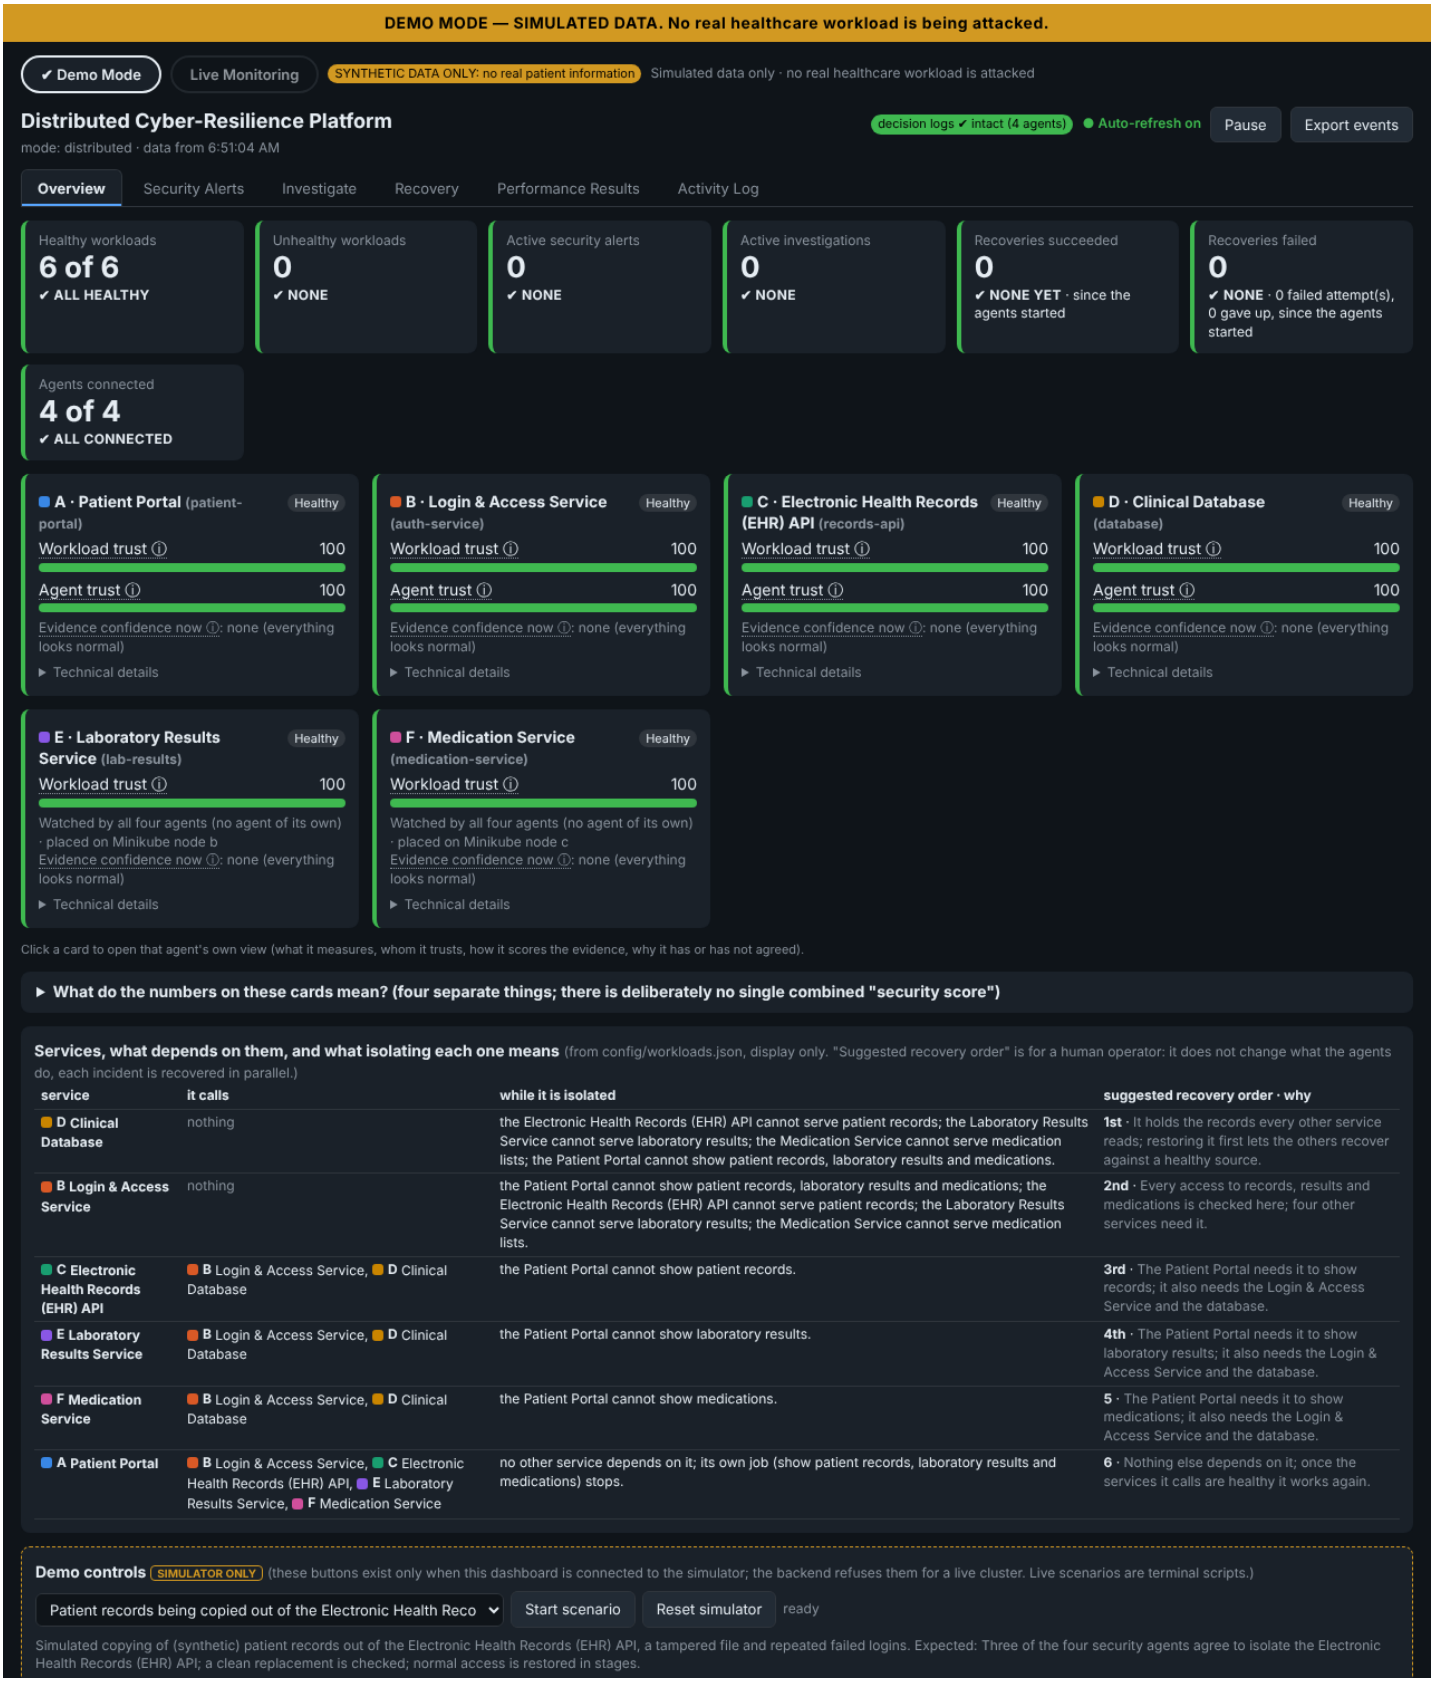
Cyber-Resilience Platform: dashboard screenshots (Phases A-I)

Taken from the real simulator behind the dashboard (Demo Mode): four agents, six healthcare workloads, synthetic data only.
Final commit ede6e28 (Phase I). Not a live cluster; the Live page is shown not connected.
Page 1 is this note; each following page is one screenshot.

Overview (Demo Mode, simulator): all six healthcare workloads, healthy



01-overview-six-workloads.png

Services, what depends on them, and what isolating each one means

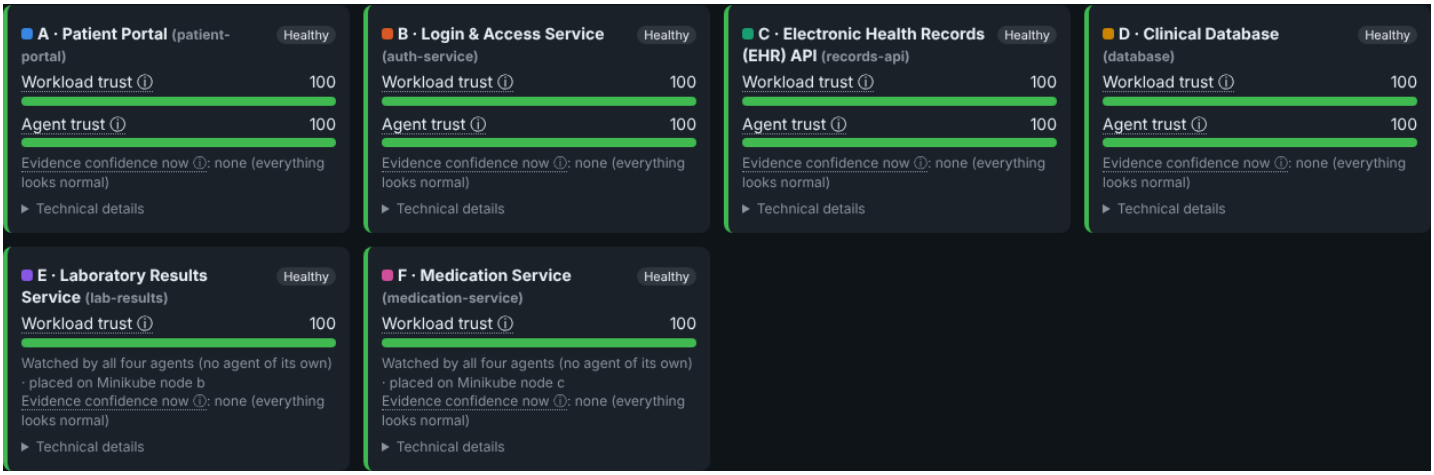


Services, what depends on them, and what isolating each one means (from config/workloads.json, display only. "Suggested recovery order" is for a human operator: it does not change what the agents do, each incident is recovered in parallel.)

service	it calls	while it is isolated	suggested recovery order · why
D Clinical Database	nothing	the Electronic Health Records (EHR) API cannot serve patient records; the Laboratory Results Service cannot serve laboratory results; the Medication Service cannot serve medication lists; the Patient Portal cannot show patient records, laboratory results and medications.	**1st** · It holds the records every other service reads; restoring it first lets the others recover against a healthy source.
B Login & Access Service	nothing	the Patient Portal cannot show patient records, laboratory results and medications; the Electronic Health Records (EHR) API cannot serve patient records; the Laboratory Results Service cannot serve laboratory results; the Medication Service cannot serve medication lists.	**2nd** · Every access to records, results and medications is checked here; four other services need it.
C Electronic Health Records (EHR) API	**B** Login & Access Service, **D** Clinical Database	the Patient Portal cannot show patient records.	**3rd** · The Patient Portal needs it to show records; it also needs the Login & Access Service and the database.
E Laboratory Results Service	**B** Login & Access Service, **D** Clinical Database	the Patient Portal cannot show laboratory results.	**4th** · The Patient Portal needs it to show laboratory results; it also needs the Login & Access Service and the database.
F Medication Service	**B** Login & Access Service, **D** Clinical Database	the Patient Portal cannot show medications.	**5** · The Patient Portal needs it to show medications; it also needs the Login & Access Service and the database.
A Patient Portal	**B** Login & Access Service, **C** Electronic Health Records (EHR) API, **E** Laboratory Results Service, **F** Medication Service	no other service depends on it; its own job (show patient records, laboratory results and medications) stops.	**6** · Nothing else depends on it; once the services it calls are healthy it works again.

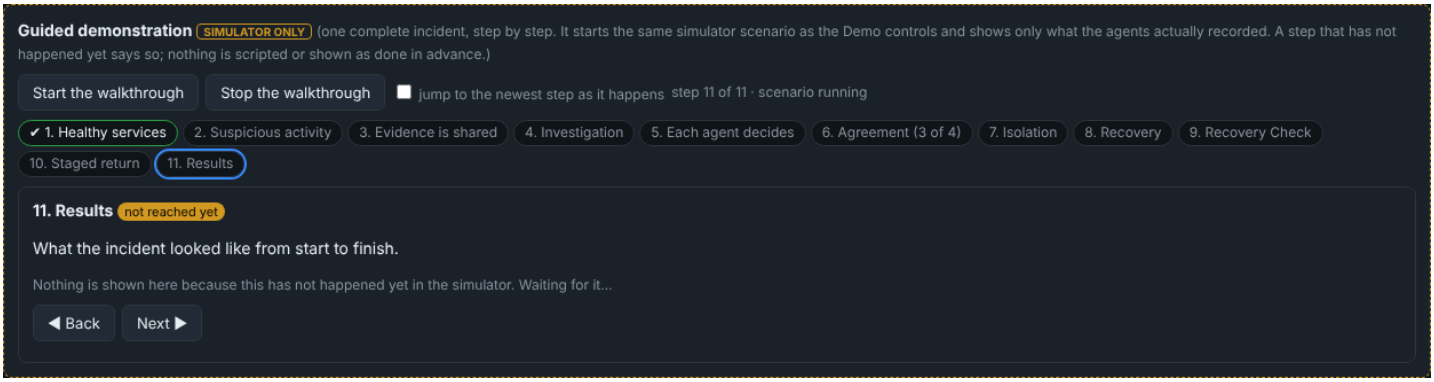
02-services-and-impact.png

Six workload cards: E and F are watched by all four agents (no agent of their own)



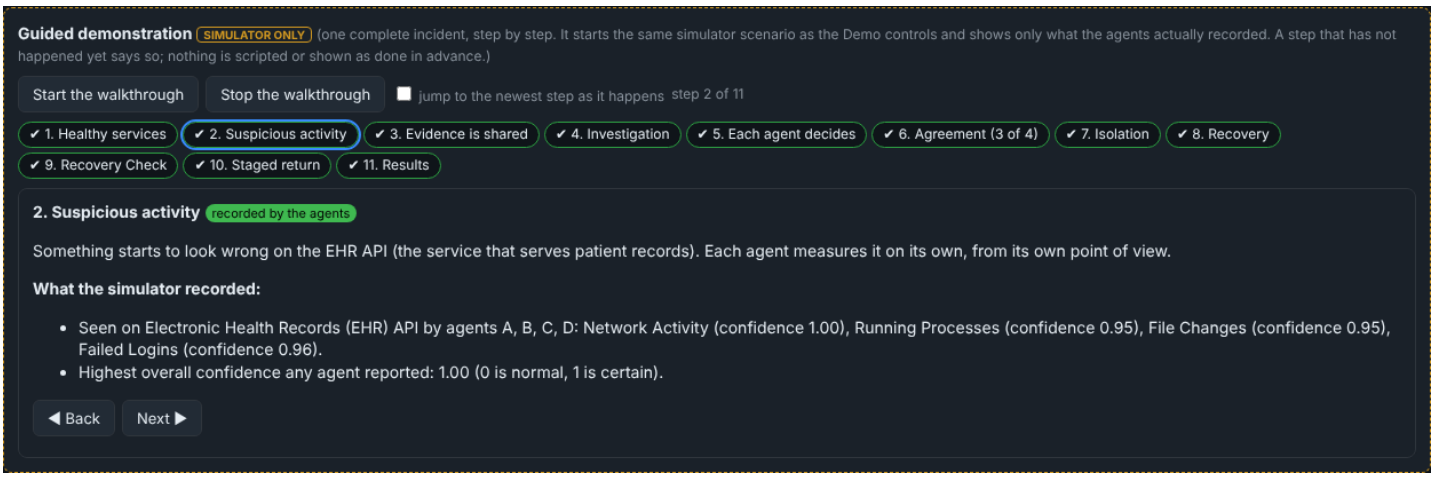
03-six-workload-cards.png

Guided demonstration: a step that has not happened yet says 'not reached yet'



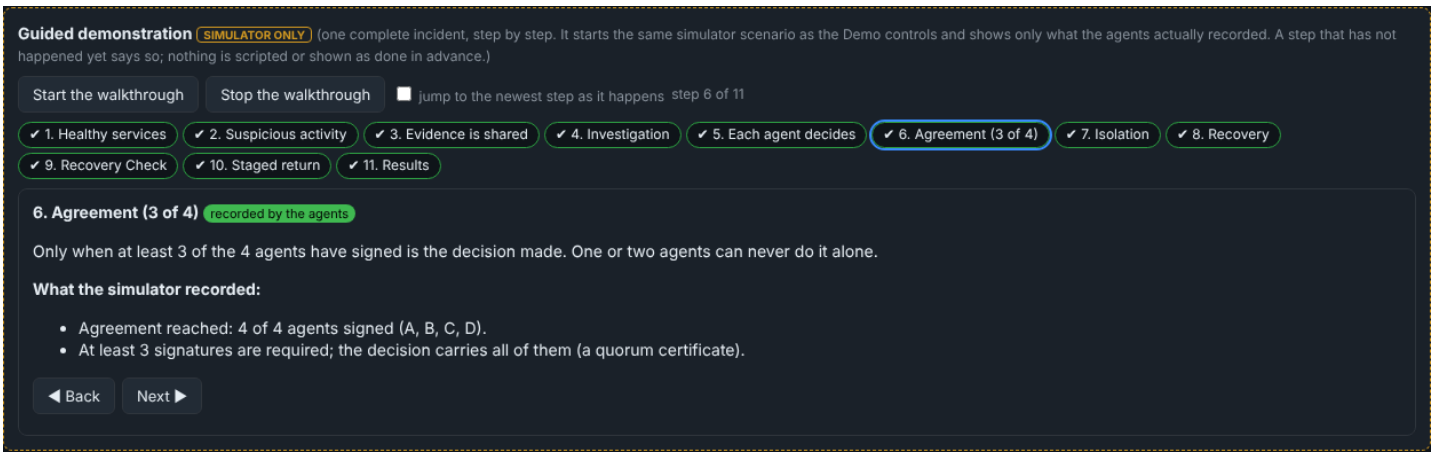
04-walkthrough-step-not-reached.png

Walkthrough step 2: suspicious activity on the EHR API, as the agents measured it



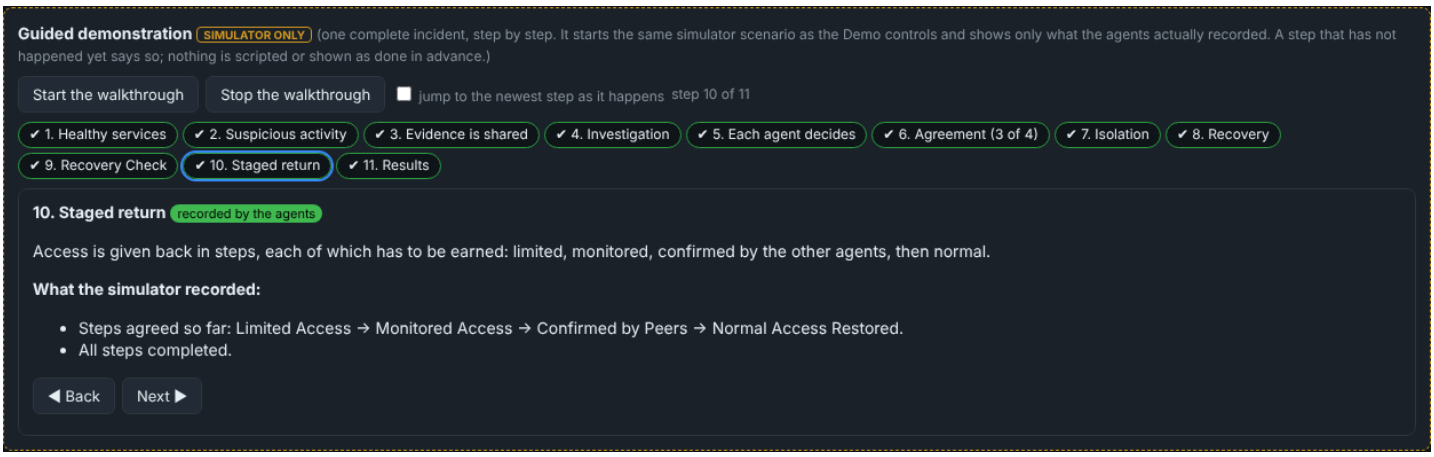
05-walkthrough-2-suspicious.png

Walkthrough step 6: agreement (3 of 4 agents signed)



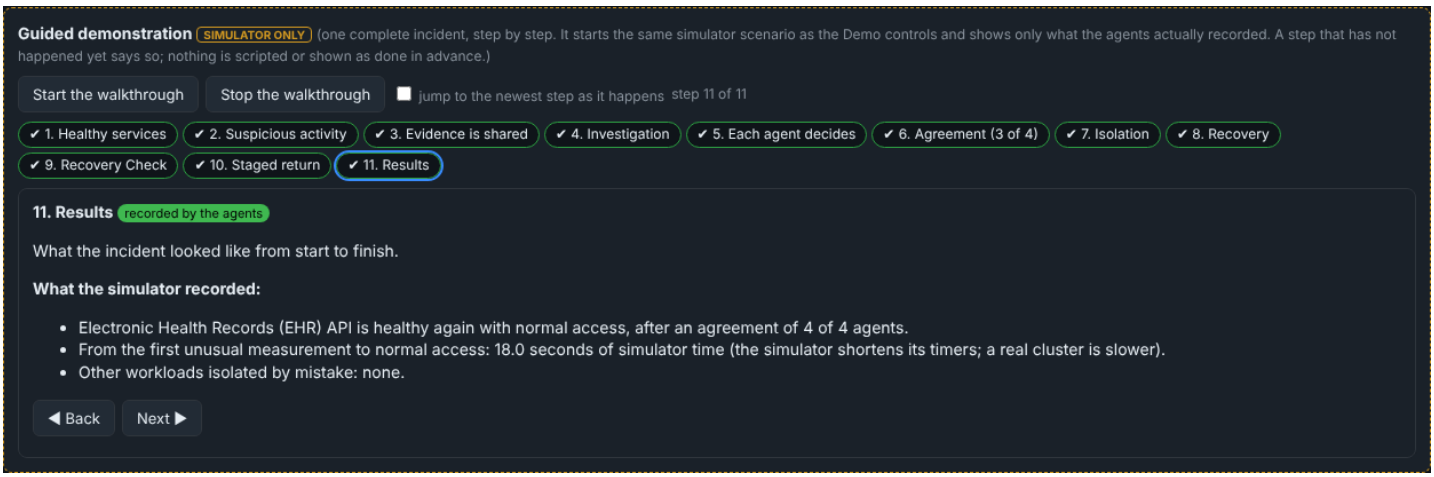
06-walkthrough-6-agreement.png

Walkthrough step 10: staged return to normal access



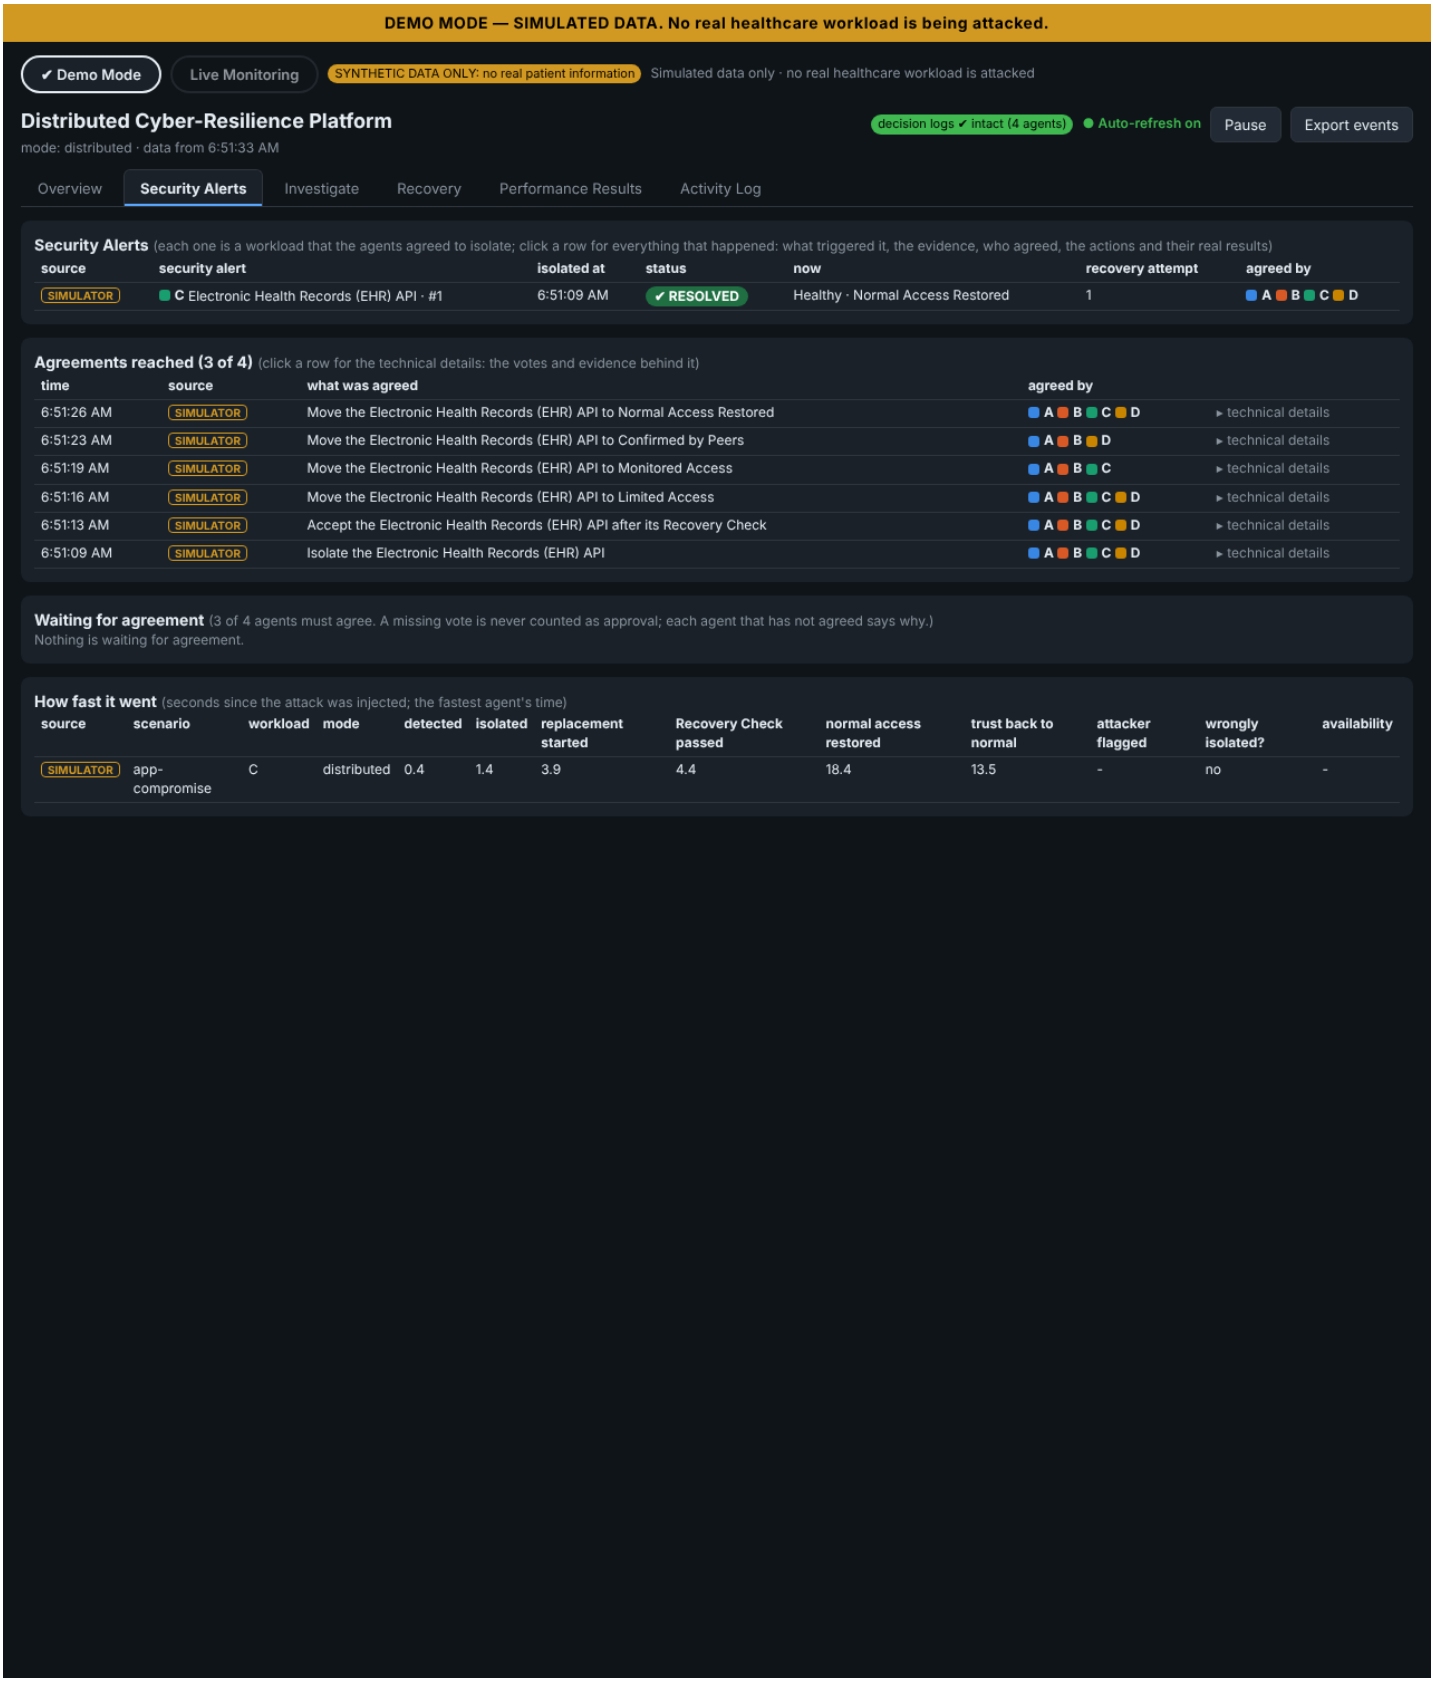
07-walkthrough-10-staged-return.png

Walkthrough step 11: results (simulator time)



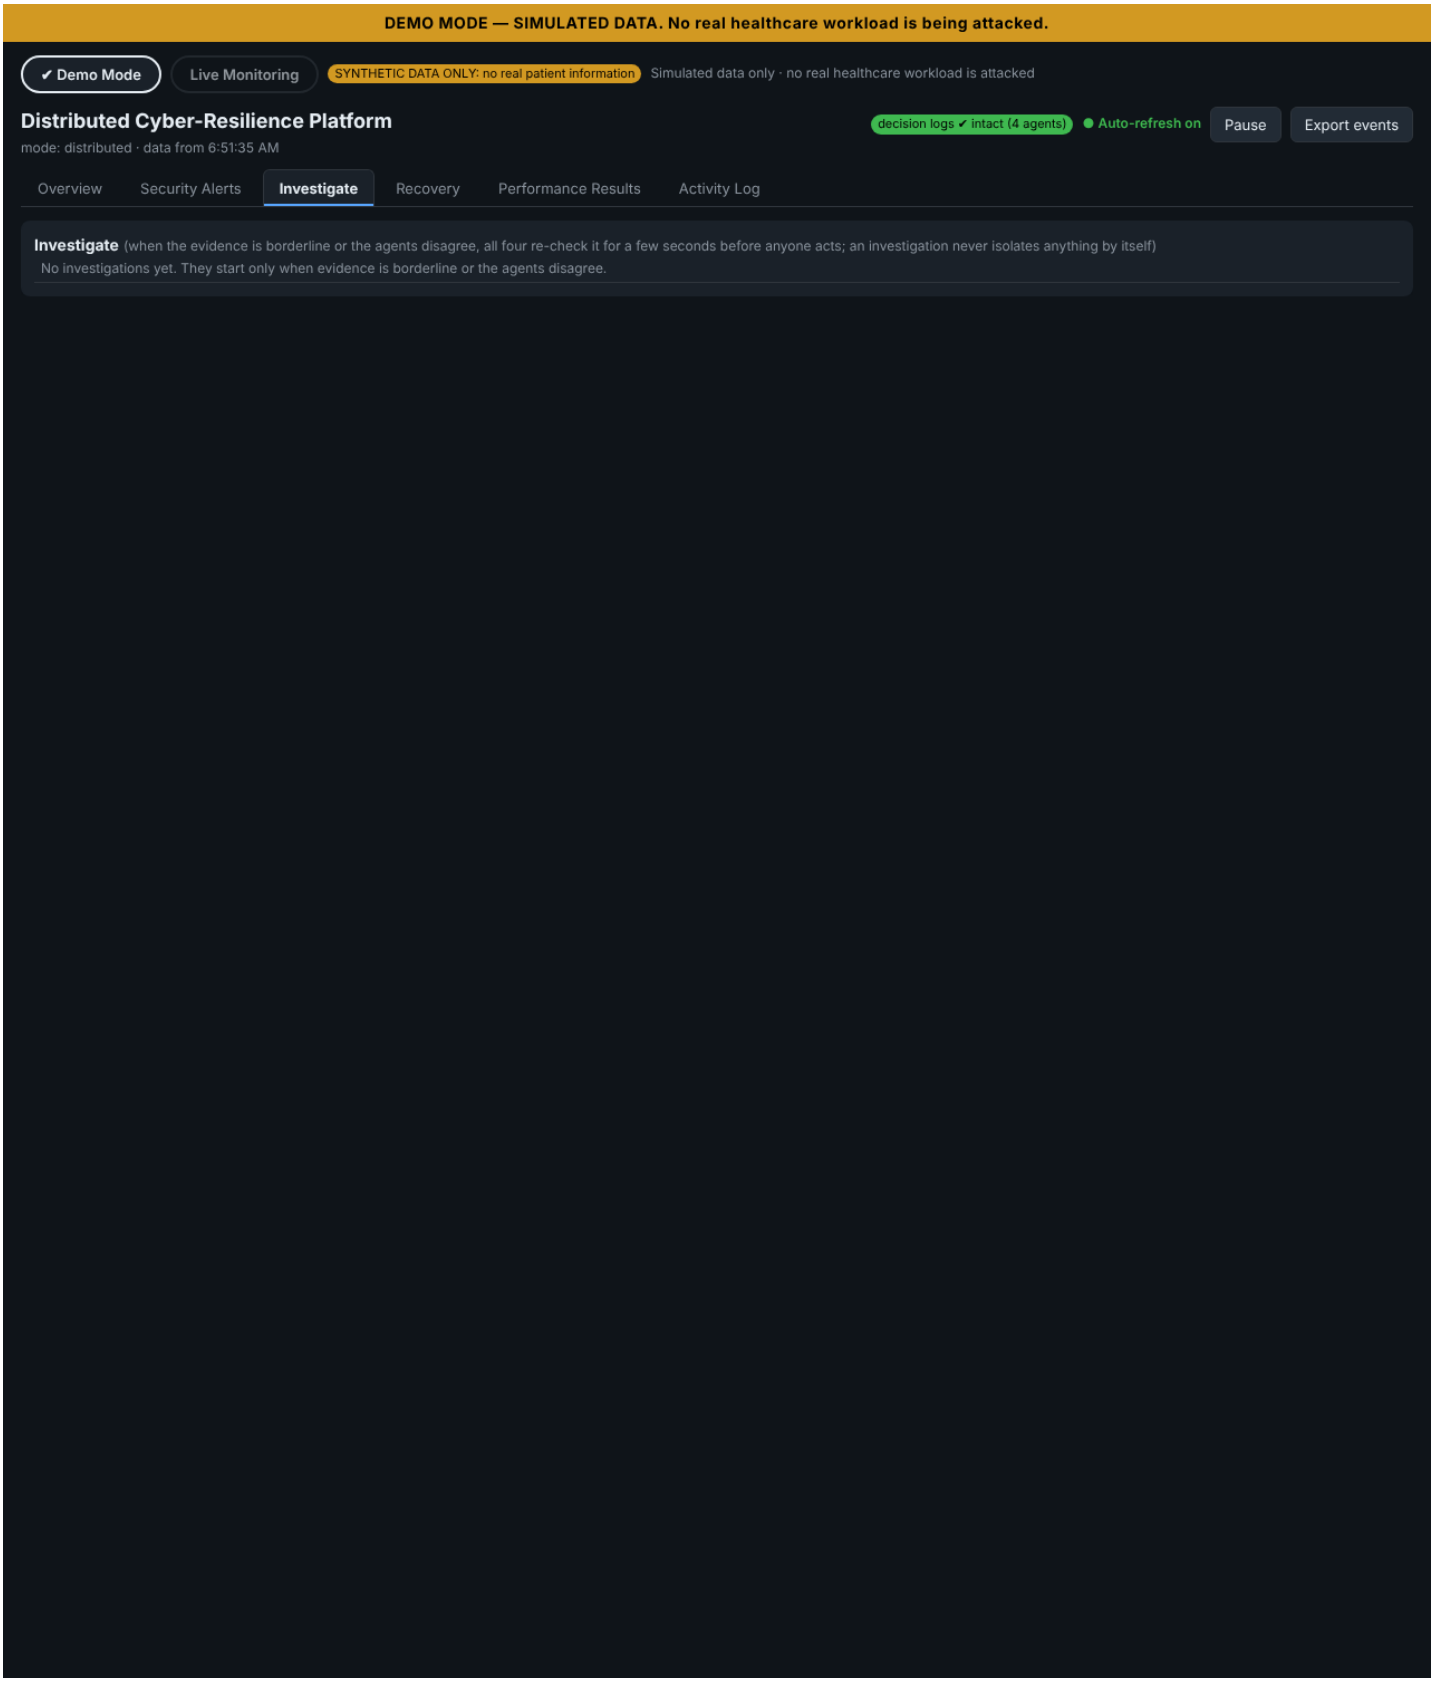
08-walkthrough-11-results.png

Security Alerts tab: the recorded incident



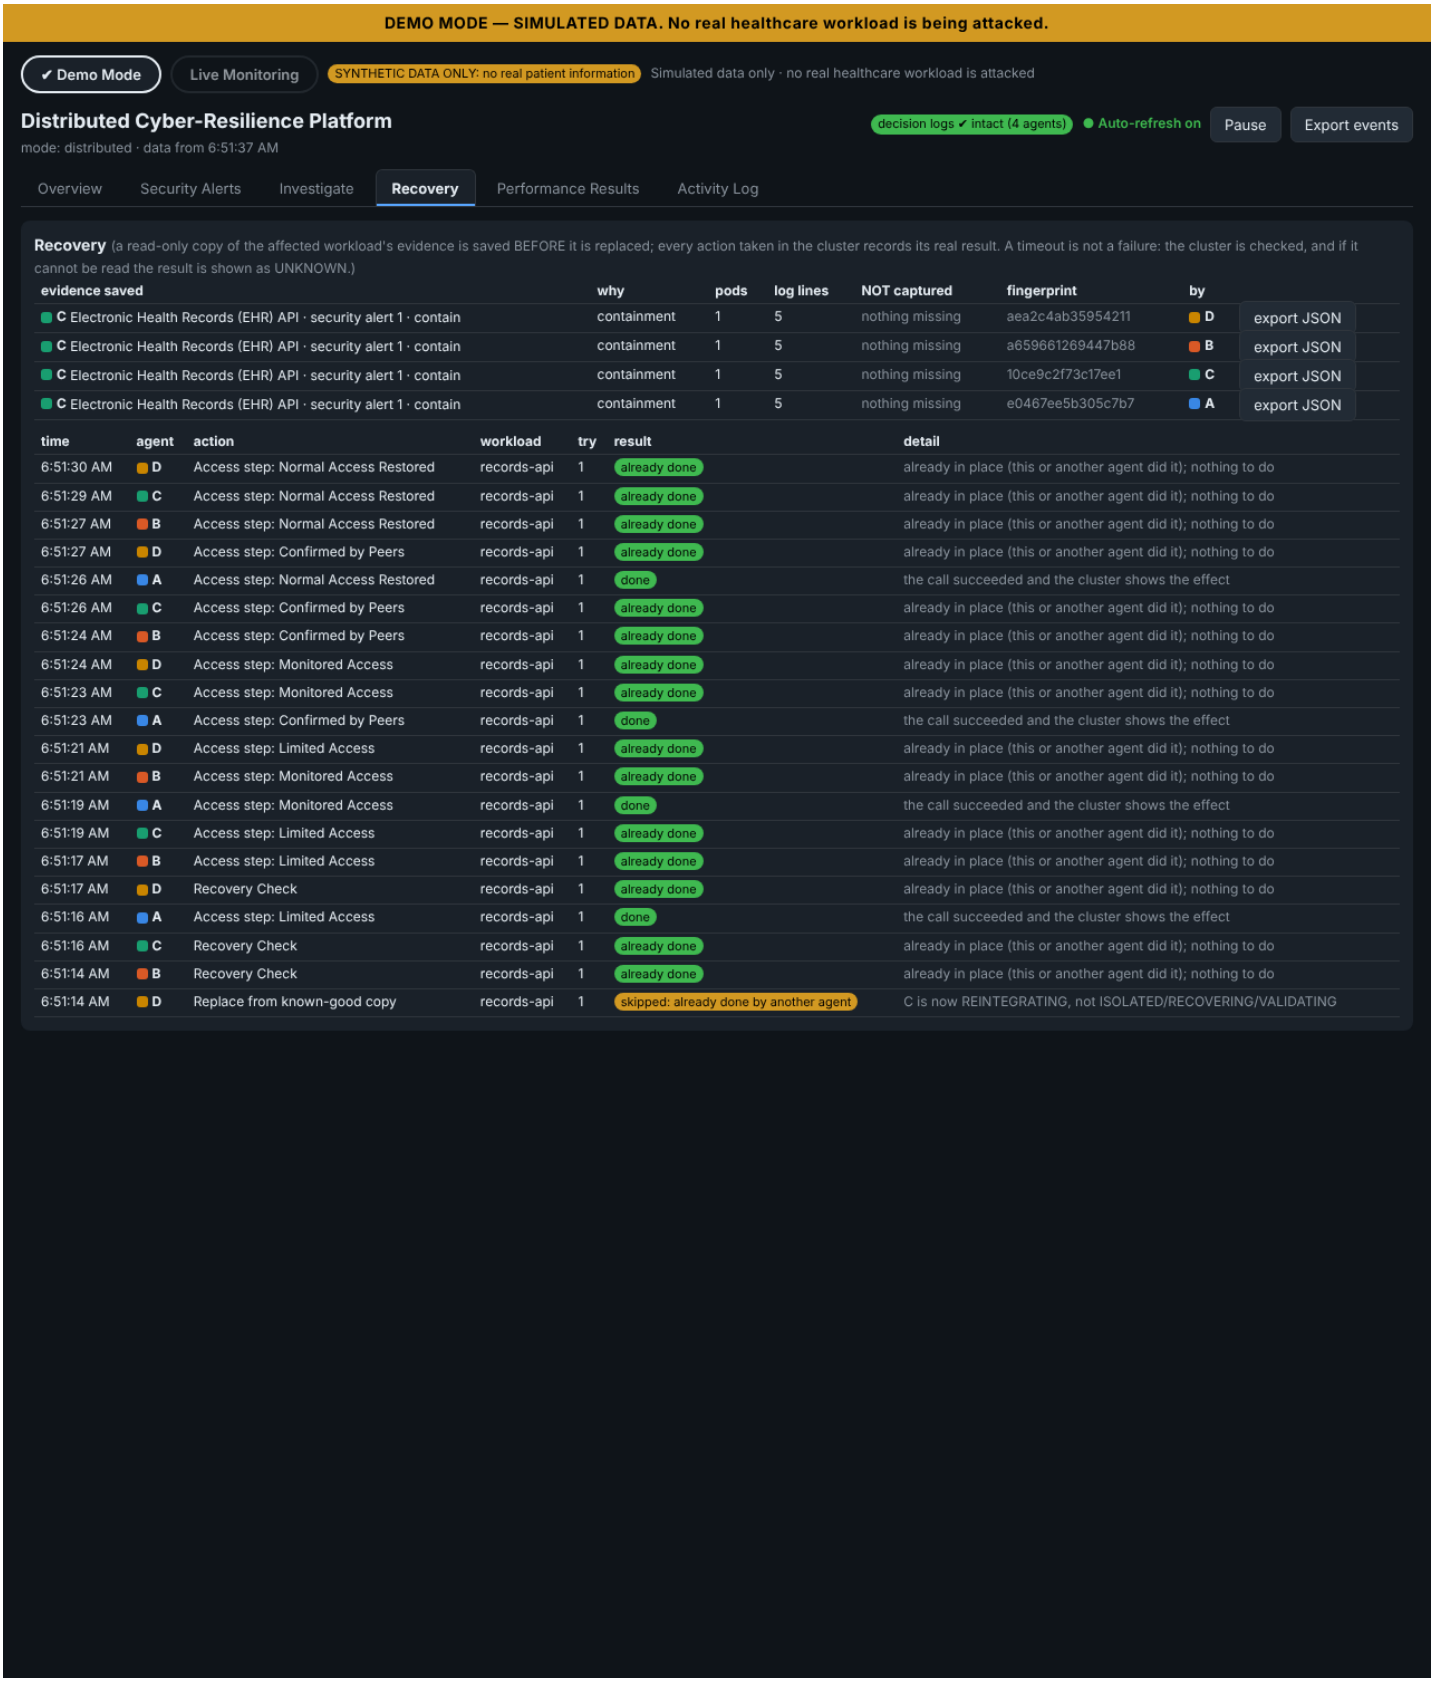
09-security-alerts.png

Investigate tab



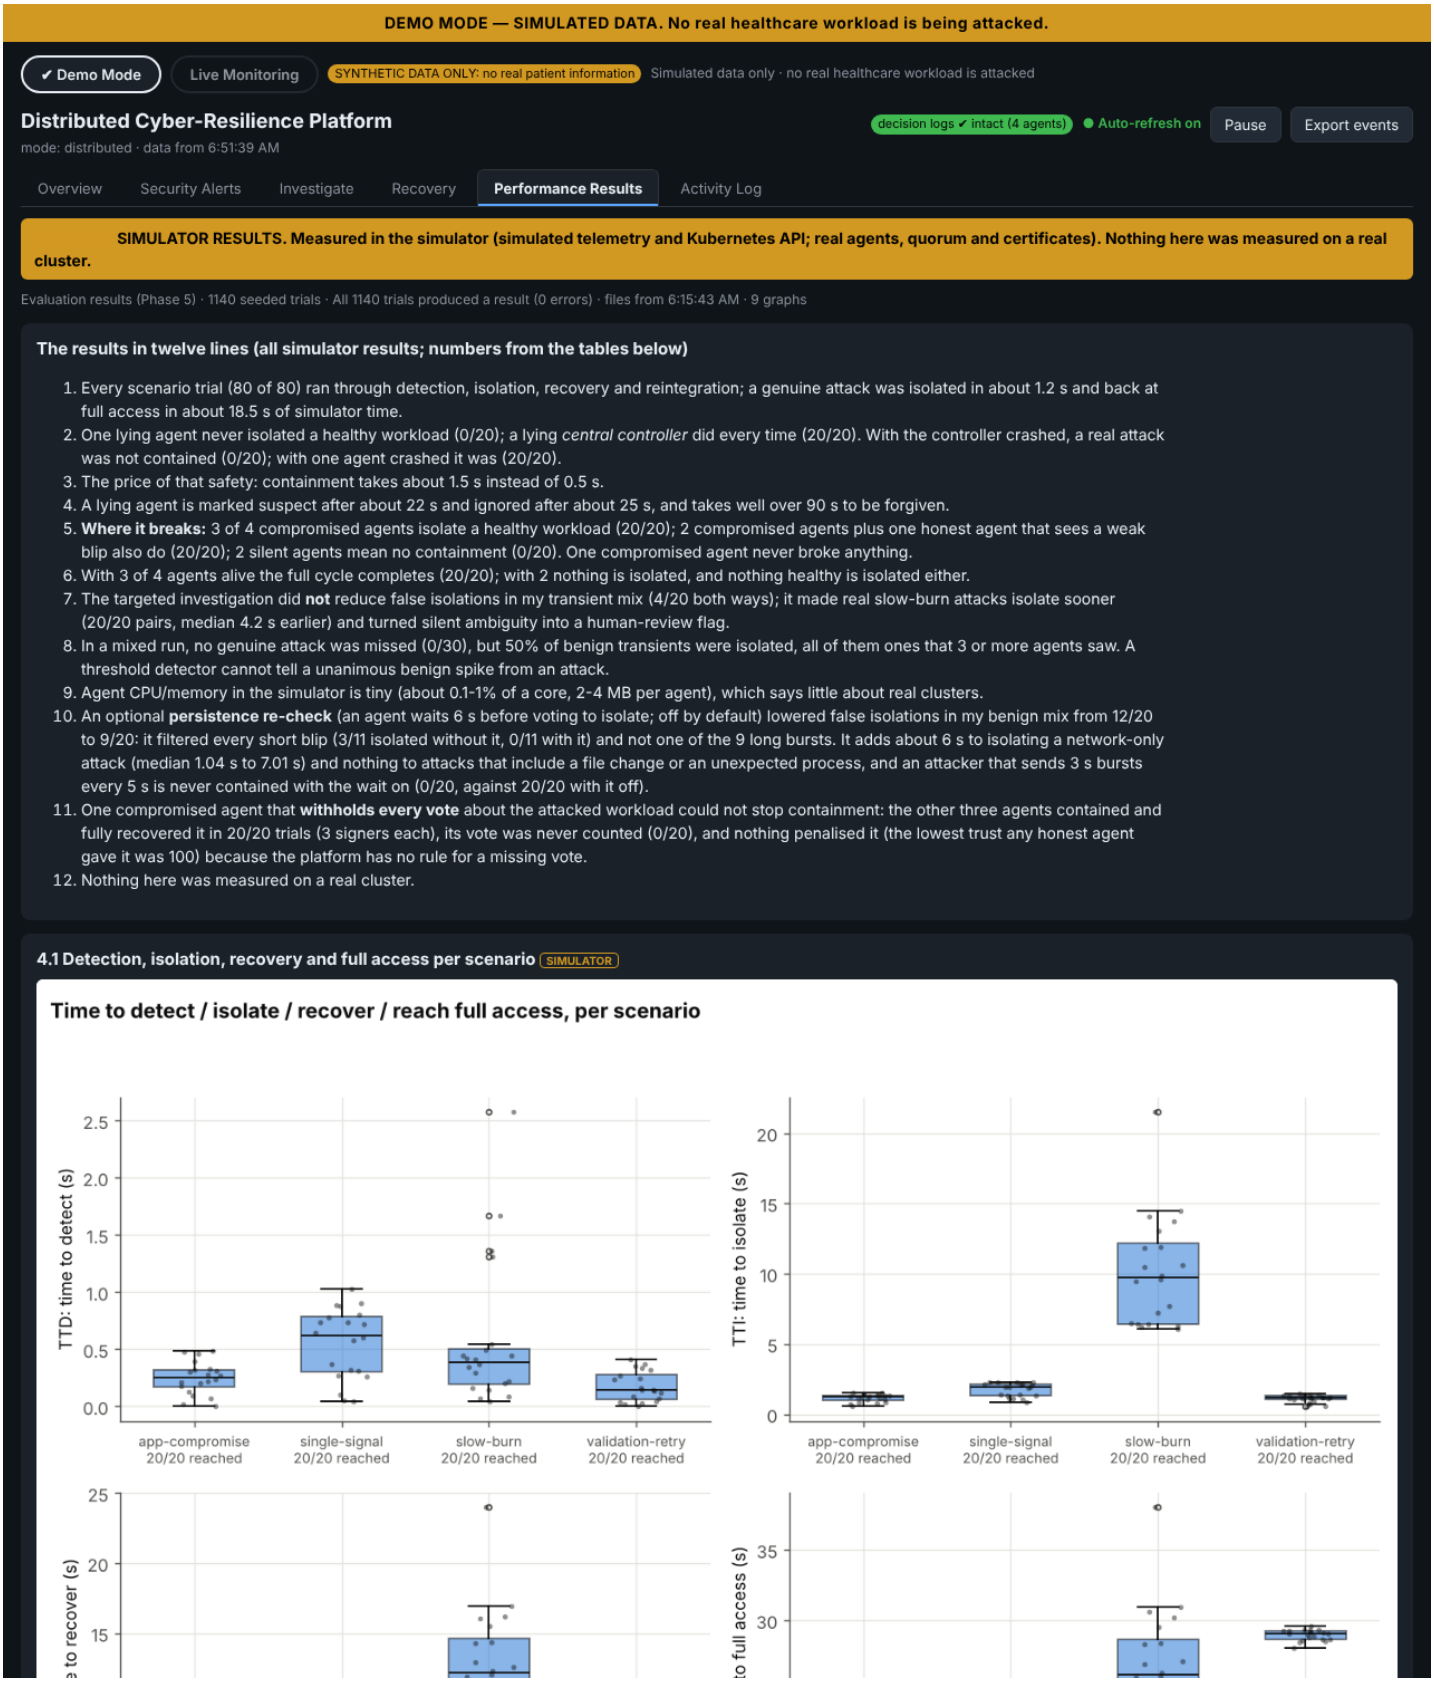
10-investigate.png

Recovery tab



11-recovery.png

Performance Results tab (simulator results)

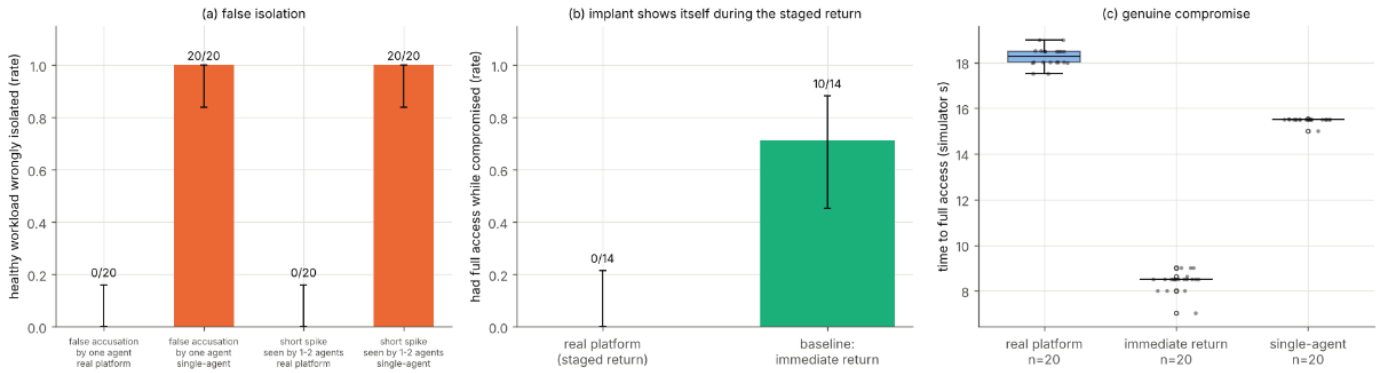


12-performance-results-top.png

Performance Results 4.11 (Phase I): evaluation-only baselines vs the real platform

4.11 What the safeguards buy: two evaluation-only baselines (Phase I) SIMULATOR

Evaluation-only baselines against the real platform (simulator, identical seeds)



Simulator results (simulated telemetry and Kubernetes API; real agents, quorum and certificates)

Metric	Start event	End event	Denominator
Contained	attack injected	any agent isolates the workload within 40 s	all trials of the row
Time to isolate	attack injected	first agent isolates it (same as TTI above)	trials that isolated
Time to full access	attack injected	first agent confirms normal access (TTF)	trials that reached it (the table says n)
Full access while compromised	the dormant implant starts (injected after the stated delay, once the replacement began its return)	any 0.25 s sample in which the workload has **no isolation policy** (full access) **and** is compromised	trials in which the replacement reached the return stage (the "returning" trials; the table gives the number)
Seconds at full access while compromised	same start	sum of those samples x 0.25 s	same trials
Contained again / ended clean	the implant starts	a second isolation (epoch 2) / back at full access with a clean replacement	same trials
Healthy workload wrongly isolated	accusation signed / spike started	any agent isolates a workload that was never attacked within 25 s / 30 s	all trials of the row

Policy	Trials (denominator)	Contained	Time to isolate, simulator s	Time to full access, simulator s	Reached full access	Another workload isolated
real platform (3 of 4, staged return)	20	20/20 = 100% (95% CI 84-100%)	median 1.1 (IQR 1.1-1.1, min 1.1, max 1.5, n=20)	median 18.3 (IQR 18.0-18.5, min 17.5, max 19.0, n=20)	20/20 = 100% (95% CI 84-100%)	0/20
baseline 1: immediate reintegration (no stage gates)	20	20/20 = 100% (95% CI 84-100%)	median 1.1 (IQR 1.1-1.1, min 0.6, max 1.6, n=20)	median 8.5 (IQR 8.5-8.5, min 7.0, max 9.0, n=20)	20/20 = 100% (95% CI 84-100%)	0/20
baseline 2: single-agent decision (quorum 1)	20	20/20 = 100% (95% CI 84-100%)	median 0.5 (IQR 0.5-0.5, min 0.0, max 0.5, n=20)	median 15.5 (IQR 15.5-15.5, min 15.0, max 15.5, n=20)	20/20 = 100% (95% CI 84-100%)	0/20

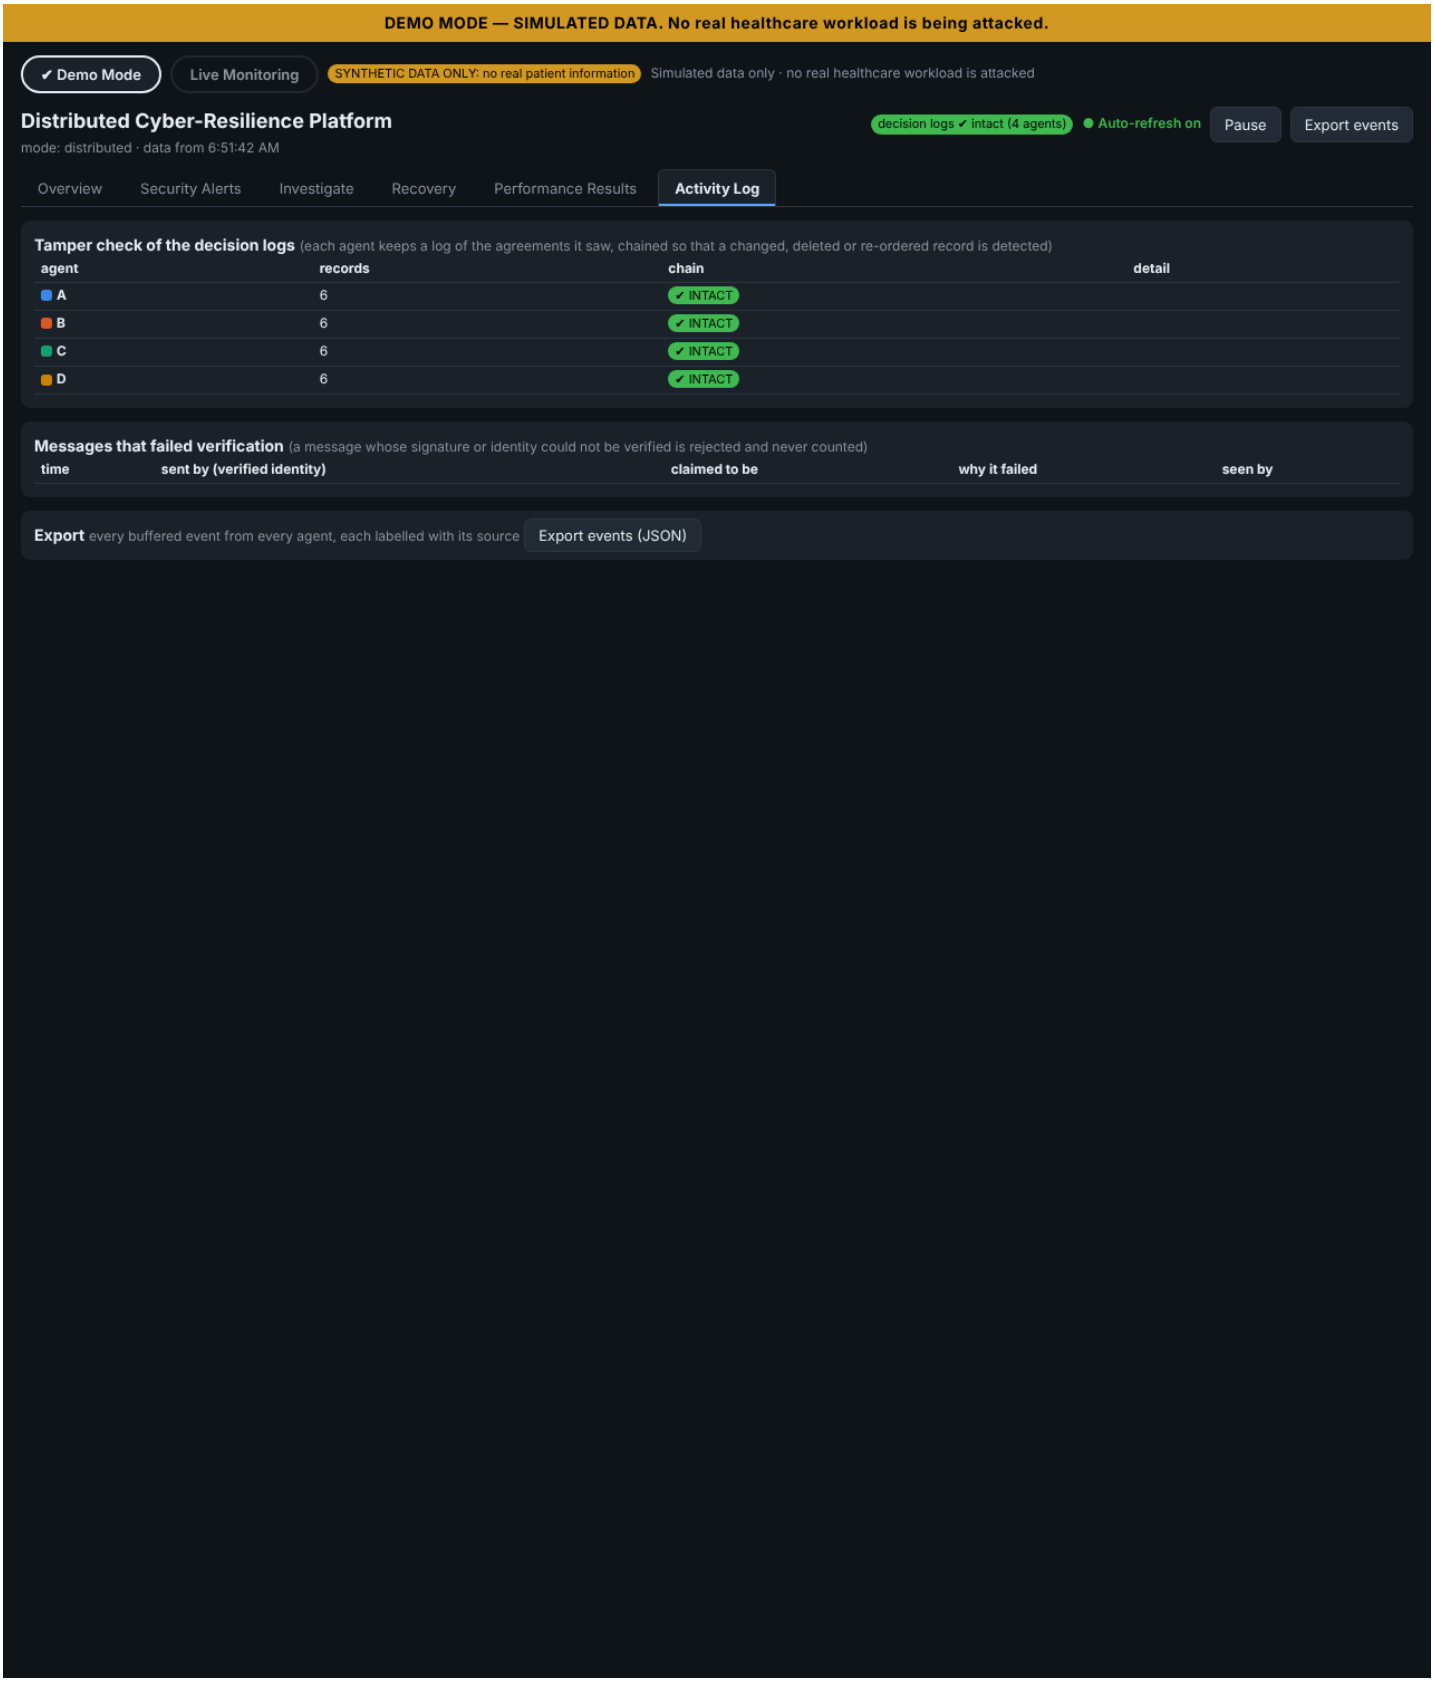
Implant shows itself	Policy	Trials that reached the return stage (denominator)	Full access while compromised	Seconds at full access while compromised	Contained again	Ended clean
during the staged return (2-12 s after the return began)	real platform (3 of 4, staged return)	14	0/14 = 0% (95% CI 0-22%)	median 0.00 (IQR 0.00-0.00, min 0.00, max 0.00, n=14)	14/14 = 100% (95% CI 78-100%)	14/14 = 100% (95% CI 78-100%)
during the staged return (2-12 s after the return began)	baseline 1: immediate reintegration (no stage gates)	14	10/14 = 71% (95% CI 45-88%)	median 1.00 (IQR 0.19-1.00, min 0.00, max 1.50, n=14)	14/14 = 100% (95% CI 78-100%)	14/14 = 100% (95% CI 78-100%)
after the staged return would have finished (> 12 s)	real platform (3 of 4, staged return)	6	5/6 = 83% (95% CI 44-97%)	median 1.00 (IQR 0.81-1.00, min 0.00, max 1.00, n=6)	6/6 = 100% (95% CI 61-100%)	6/6 = 100% (95% CI 61-100%)
after the staged return would have finished (> 12 s)	baseline 1: immediate reintegration (no stage gates)	6	6/6 = 100% (95% CI 61-100%)	median 1.00 (IQR 0.81-1.19, min 0.75, max 1.50, n=6)	6/6 = 100% (95% CI 61-100%)	6/6 = 100% (95% CI 61-100%)

Situation	Policy	Trials (denominator)	Healthy workload wrongly isolated	Time to the wrong isolation, simulator s
one compromised agent signs a false accusation	real platform (3 of 4, staged return)	20	0/20 = 0% (95% CI 0-16%)	n/a (no trial reached it)
one compromised agent signs a false accusation	baseline 2: single-agent decision (quorum 1)	20	20/20 = 100% (95% CI 84-100%)	median 0.8 (IQR 0.8-0.8, min 0.2, max 0.8, n=20)
a short traffic spike seen by 1-2 agents	real platform (3 of 4, staged return)	20	0/20 = 0% (95% CI 0-16%)	n/a (no trial reached it)
a short traffic spike seen by 1-2 agents	baseline 2: single-agent decision (quorum 1)	20	20/20 = 100% (95% CI 84-100%)	median 0.8 (IQR 0.6-0.8, min 0.2, max 0.8, n=20)

▶ **Explanation: how to read it and what the data shows**

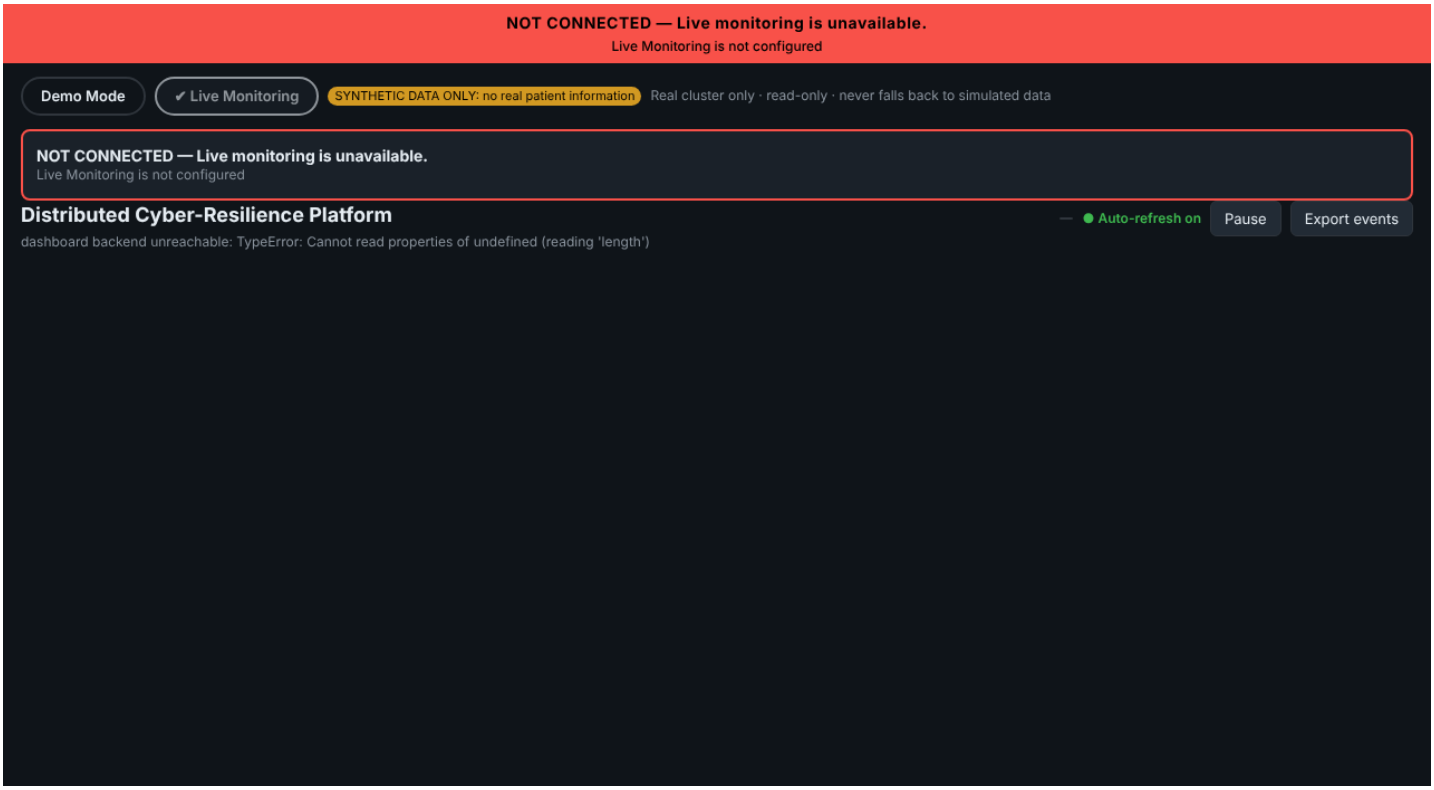
13-baselines-section-4-11.png

Activity Log tab



14-activity-log.png

Live Monitoring page: honestly 'not connected' here (no real cluster); never shows simulated data



15-live-page-not-connected.png